The Control: a follow-up study
After Weekly Comp #14 · run Thursday 15 October 2026

Does an ICM workspace beat the folder?

The competition entry tested a five-file folder against the same documents with no instructions. This follow-up adds a fourth set-up, an ICM workspace built by the ICM method, and runs all four on the same 18 documents, the same model and the same rule.

> The ICM was the most accurate set-up: right on all 54 answers by both markers. It also read the most. By the rule fixed before the run, it did not earn its place.

The four set-ups

Same 18 documents in every one, same model (Claude Sonnet 5.5), same tools (read and search only), same questions, run in one batch in random order.

🟦 Documents only
The 18 documents and nothing else.
0
characters of instructions loaded first

🟧 The entry's folder
A map, identity, rules and examples, all loaded before the question.
5,261
characters loaded first

🟩 ICM workspace
A short router loads first. The stage contract and document list are read when needed.
902
characters loaded first

🟪 Same ICM, all at once
The ICM's files, word for word, but every one loaded before the question.
11,886
characters loaded first

The results

18 new questions, 3 runs each: 54 answers per set-up. Means per question unless it says otherwise. Shaded cells are the best in each row.

	🟦 Documents only	🟧 Folder	🟩 ICM	🟪 ICM, all at once
RIGHT ANSWERS, OUT OF 54				
Marked by Claude Opus 5.5 arm hidden	52	52	**54**	**54**
Marked by Gemini 4 Flash another maker, arm hidden	50	52	**54**	**54**
Hold every fact the key asks for counted by script, no marker	49	50	53	**54**
WORK				
Steps to find the documents searches, listings, document opens	5.44	**2.46**	3.61	3.89
Instruction files opened	**0**	**0**	2.00	1.96
Seconds	7.8	**4.0**	9.0	8.7
Words in the answer	166	**85**	312	318
COST				
Text put in front of the model characters, loaded and fetched	**8,171**	10,447	17,781	29,208
Tokens in whole conversation, every turn	22,711	**18,372**	26,876	28,090
Tokens out	925	**443**	1,176	1,203
Cost per question list price, US dollars	0.026	**0.014**	0.027	0.034

Tokens in count the whole conversation each time it is sent, so they grow with the number of turns as well as with the text; the documents-only set-up took about six turns a question. That is why the rule uses characters of text, not tokens.

Right answers, out of 54

Three readings of the same answers. Higher is better.

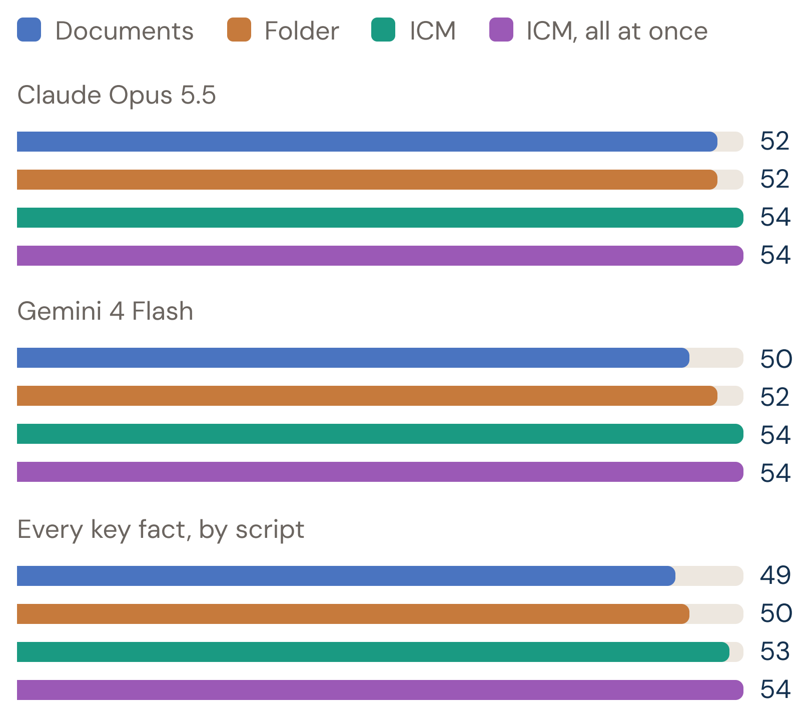

What each answer cost

Per question. Shorter is cheaper.

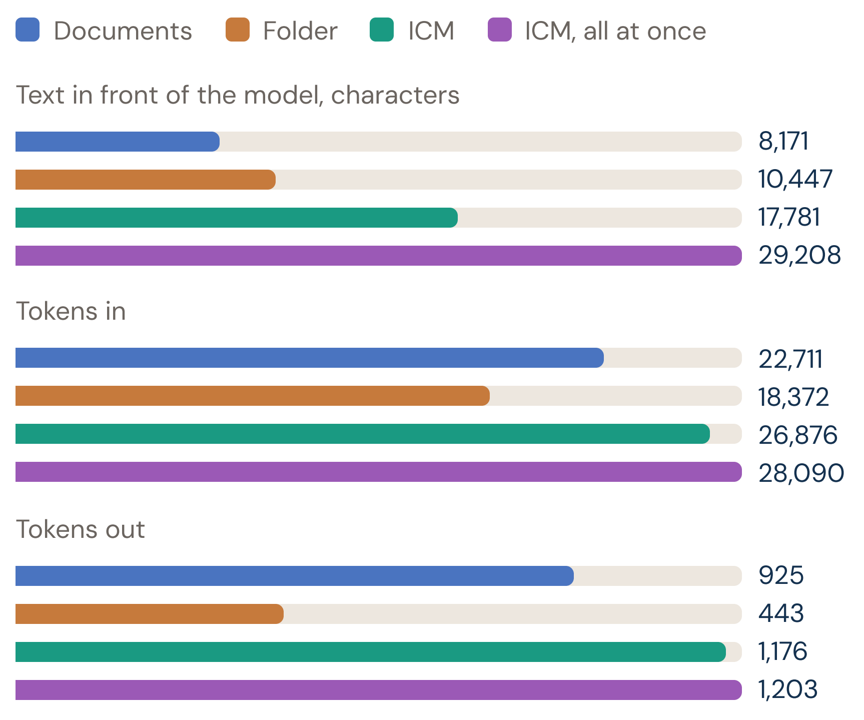

The verdicts, by the rule fixed before the run

The same three-part rule as the competition entry: a set-up earns its place only if it finds the documents in clearly fewer steps, loses no right answers, and puts no more than 5% more text in front of the model.

Did not earn its place **ICM against the documents alone**	Fewer steps on 16 of 18 questions and no answers lost. It failed on text alone: 17,781 characters against 8,171.
Did not earn its place **ICM against the folder**	More steps to find the documents (more on 14 of 18 questions) and more text: 17,781 characters against 10,447.
Could not tell **Reading on demand against loading it all first**	The same ICM files, loaded when needed instead of all at once, saved 39% of the text and changed nothing else: same answers, about the same steps. The all-at-once version still opened its two instruction files in almost every session, although they were already loaded.

The rule only checks that a set-up does not lose right answers; it gives no credit for gaining them. The ICM's extra accuracy does not count toward earning its place.

What it means

- **The ICM bought accuracy with reading.** Its contract tells the model to read every paper that touches the question, check every change, and show its working before answering. It got every answer right, and it wrote answers almost four times as long as the folder's.
- **The folder is still the cheapest way to the documents.** Its map took the model to the right papers in 2.46 steps, and the entry's own result came back in this batch: fewer steps on 10 of 10 of the entry's questions in both sets, with more text and fewer complete answers than the documents alone.
- **On-demand loading is real, and small here.** At 18 documents the saving was text, not behaviour. The method places its saving against libraries many times this size.
- **Most questions did not separate the set-ups.** On 12 of 18, every set-up was right every time. The differences come from six questions.
- **How it fits the ICM paper.** Jake Van Clief and David McDermott's paper says that no controlled comparison of staged loading against loading everything at once has been done yet. This is a small first one, at a scale where the paper would not expect a quality difference, and there was none. The paper's case is at 30,000 to 50,000 tokens.

How it was kept fair

Questions written first
18 new questions written by a separate Claude Opus 5.5 session from the documents alone, key-checked by another model, committed before the ICM existed. I never read them until the run was over.

Practice matched
The ICM got the same five practice rounds the folder had, then had to pass two gate audits by sessions that did not build it.

Markers blind, two makers
Answers shuffled under random numbers. Four planted answers of known grade were marked right by both. Neither marker could guess the set-up (44% and 32%).

The verdict is a script
The rule, the thresholds and the sentences were code, tested on made-up data before the run. It was run once.

The ICM built blind
Built by the icm-architect method on practice documents only, in the house style of my PO Verifier. The builder never saw the live documents or questions.

Plan reviewed four times
Three design reviews (method, ICM fidelity, a skeptic), then one by Perplexity, which added the all-at-once set-up. Every change is logged.

One sealed run, every session kept
456 sessions, no faults. A first run was voided when a usage limit refused a quarter of its sessions; it is published, never analysed.

Predictions on record
Written before the run: fewer steps than the documents alone, and more text than both. Both held.

What this cannot tell you

- Anything about ICM's hand-offs between stages or its human checks. For a single question the method chose one stage: there is nothing to hand on and nobody to check it.
- Anything about large libraries, where ICM claims its biggest saving. This one is about 32,000 characters.
- Anything about every possible ICM. This one passed every check and two audits, but its one stage contract is about three times the size the method recommends. Its instruction to read every relevant paper and show the working most likely bought both the extra accuracy and the extra cost.
- Anything beyond one model and made-up rule documents. The new questions were written and checked by Claude models.

Everything is public: the plan, every review, the ICM, the questions, all 912 session logs, both markers' marks and the scripts.
github.com/jdvanleenen-sys/the-control-icm · Start with RESULTS.md. More tests to follow.